

WHITE PAPER

V2026.2 — DATED ON PURPOSE

The *Dual-Bounded* Board

How AI is already reshaping strategic judgment — and what governs it next.

THE-PARADOX.COM · V2026.2 · CURRENT AS OF THE CAPABILITY FRONTIER DESCRIBED BELOW

Your board has already delegated part of its judgment to artificial intelligence. It did not vote on it, minute it, or govern it, and most directors could not tell you where it happened.

AI-influenced material now enters strategic decisions through informal channels: a paper shaped by a model, an option set quietly pre-narrowed before it reaches the table, an analysis no one flagged as machine-assisted. Across sixteen board-level interviews conducted for the research behind this paper, this was not the exception. It was the condition.

So the opening question is not whether to integrate AI into board decisions. Integration has already occurred. The question is whether it is governed, and in strategic decisions under genuine uncertainty, where there is no objectively correct answer and the quality of judgment is the entire game, an ungoverned influence on that judgment is not a technical footnote. It sits precisely where a board is supposed to earn its keep.

| *Every board is now a human–AI decision system. Most do not know it yet.*

Three conversations, one gap

Three conversations are converging on this gap, and it helps to be clear about how they relate.

In *Harvard Business Review*, Jonathan Rosenthal and Neal Zuckerman argue that consensus decision-making fails in the AI era, and push boards toward sharper decision rights and structured disagreement.¹ That is a prescription: change how the decision is owned.

In *Co-Existence*, Ethan Mollick maps what he calls the jagged frontier, the strange in-between years in which AI is sometimes, but not always, sharper than the people beside it, and in which the cooperative model of the early chatbot era is already giving way to more autonomous systems.² That is a description of the terrain.

This paper occupies the ground between them. Not a prescription asserted from first principles, nor the terrain described from the individual's vantage point, but an empirical account of how a board actually behaves as a single system operating on that frontier, drawn from sixteen interviews and 1,080 coded passages with European decision-makers at board level. This research documented thirteen paradoxes of human-AI collaborative decision-making.³ Three are examined here, those that boards encounter first.

The prescription tells a board how to decide. The terrain tells it what it is standing on. What has been missing is the part in between: what actually happens to a board's judgment once AI is in the room, and what to do about it. That is the subject here.

One system, not two

Everything that follows depends on one shift in how the board sees itself.

The instinctive model is two parties: a human decision-maker consulting a tool and retaining the final say. That model is wrong, and the error is expensive. A board and its AI form a single bounded system.

Human judgment is bounded in the ways it has always been. We satisfice rather than optimise. We anchor on the first credible answer. We narrow our options early and defend the narrowing afterwards. AI is bounded too, but differently, fluent, fast, confident, and wrong in ways that are difficult to see from the chair, because the wrongness arrives in clean prose rather than visible hesitation.

Place these two kinds of boundedness together and their limits do not cancel. They compound. A board hopes its AI will act as a mirror, showing it more than it could see alone. Often it acts as a loudspeaker instead, returning the board's existing conviction in more articulate and better-evidenced form, and the board mistakes fluency for confirmation, and volume for validation.




FIGURE 1 — THE DUAL-BOUNDED SYSTEM · SOURCE: THE-PARADOX.COM/RESEARCH

This is the dual-bounded system. It is structural, because it arises from the meeting of two kinds of limitation, not from any single deficiency a software update could fix.

P THE PARADOX

03 – WHY AWARENESS ISN'T ENOUGH

The competitive advantage in AI is not capability. It is knowing where human and AI judgment both break down, and designing for it.

Why awareness isn't enough

The defining feature of the dual-bounded system is a gap between knowing and doing. Boards are, on the whole, aware of the risks of AI. That awareness rarely converts into changed practice, and where governance does exist, it tends to address the infrastructure rather than the influence — how AI actually shapes the judgment in the room. Three patterns from the research show how the gap operates.

01 / **PREMATURE
NARROWING**

The moment an AI returns a well-formed answer, the board's cognitive commitment to it

hardens, often before anyone has interrogated it. Awareness of bias does not prevent this — the commitment forms too early, and too quietly.

02 / **THE
EFFICIENCY
ILLUSION**

AI plainly saves time. But the time saved is typically absorbed by acceleration rather than reinvested in deeper analysis — the hours it frees up rarely return as scrutiny; they are spent reaching the next decision sooner.

03 / **MISPLACED
TRUST**

Trust in AI calibrates to its fluency, not its reliability. Worse, the boundary between what the board defers to AI and what it reserves for itself shifts over time without anyone deciding that it should.

What good looks like

If the dual-bounded system cannot be willed away, it can be governed. The research translates its findings into an architecture of four domains. The point of naming them as domains, each with an owner, is to put the question of who is responsible where it belongs – on a named person, not on "the board" in the abstract.

The architecture is sequential to build but simultaneous to live in. An organisation establishes the domains in order, because each depends on the one before; but once established, they operate at once, as a whole.




FIGURE 2 – THE DUAL-BOUNDED BOARD GOVERNANCE ARCHITECTURE

SOURCE: TIMMER (NYENRODE, 2026) · THE PARADOX

P THE PARADOX04 – WHAT GOOD LOOKS LIKE

Governance and decision authority comes first, and it belongs to the board and its senior governance bodies. Before expanding AI use, a board needs to know where AI-influenced material already enters its decisions, and to establish the governance before the use grows past it. In practice this means four things: that leadership has the AI literacy to oversee what it is governing; that decision processes carry explicit roles and decision rights; that human accountability is preserved, with verification as the mechanism that makes accountability real rather than nominal; and that business, ethics, and risk are held in a single governance frame rather than three disconnected ones.

Strategic decision quality comes next, and it belongs to executive decision-makers and their teams. The leverage point is earlier than most boards expect: debiasing has to be engineered into how the question is put to the AI, not bolted on when the answer comes back, because commitment hardens the moment the output appears. That means configuring AI to generate genuine alternatives and surface disconfirming evidence rather than confirmation; applying widening techniques such as pre-mortems and cross-industry analogy before converging; keeping human judgment as the final arbiter on genuinely strategic matters; and redirecting the time AI saves back into the analysis rather than spending it on speed.

Organisational capability comes third, and it belongs to those responsible for talent, HR, executive development, line management. The research is blunt here: both overreliance and underuse erode capability, and undifferentiated adoption quietly eliminates the very advantage AI was meant to create. The competence to work well with AI is developable, and should be treated as a real criterion in selection and promotion, not a soft skill assumed to arrive on its own.

Beneath all three runs the fourth domain: **transparency and culture**, the enabling conditions without which the others do not hold. Disclosure matters, and it is not neutral – making AI involvement visible changes behaviour, which is precisely why the rules for it have to be chosen deliberately rather than left to drift. So do platform choices, which trade capability against governability and should be made on purpose and revisited as circumstances change.

The first move

Architecture can stay abstract. The way to keep it honest is to ask, for each domain, a deliberately small question: who, specifically, does what, first – not this year, but this quarter.

DOMAIN	OWNER	FIRST ACT THIS QUARTER
Governance & decision authority	The chair	Map where AI-influenced material already enters board papers – before the next cycle
Strategic decision quality	The decision sponsor	Engineer the debiasing into the prompt, not the review of the output
Organisational capability	The executive for talent	Name AI competence as a criterion in selection and promotion
Transparency & culture	The board, collectively	Decide the disclosure rule for AI-assisted papers – knowing the rule itself changes behaviour

The discipline is in the specificity. If any first act drifts back into "govern this at board level" or "raise awareness," it has stopped being an act and become a sentiment. The test of this whole framework is whether each row names a person and a thing they can do before the next board cycle. If it does, the knowing-doing gap has somewhere to close.

What moves next

A word on why this document carries a version number.

The diagnosis is durable. The dual-bounded system and its paradoxes do not expire when the models improve, because they come from the structure of human-AI collaboration rather than from any limitation of today's tools. The response, though, is calibrated to where the frontier sits now, and the frontier is moving. As Mollick argues, the cooperative model in which a human prompts and reviews is already giving way to more autonomous systems that act with less direct supervision.

That shift does not dissolve the architecture; it sharpens the stakes inside it. The efficiency illusion intensifies as AI does more between the board's checkpoints. The trust boundary moves faster and less visibly as systems take on more of the judgment. Premature narrowing happens earlier still, when the option set is assembled by an agent before a human has framed the question. The owners stay the same; the first acts change. The chair's first move under autonomous agents is not the chair's first move under prompt-and-review.

So the architecture is offered as a snapshot, dated on purpose, a description of the dual-bounded board as it stands at this point on the frontier, built to be revised as the frontier advances rather than to pretend it won't.

A living diagnosis

This paper is the diagnosis. Turning it into a board's standing practice, a repeatable way of seeing and governing the dual-bounded system, rather than a one-off read, is the work the-paradox is taking forward now, including the instruments being developed to put these principles into a board's actual decision rhythm.

For now, the single most useful thing a board can do is the smallest one in the table above: find out where AI is already shaping its judgment, before deciding anything about where it should. Everything else follows from knowing that.

NOTES

¹ Jonathan Rosenthal and Neal Zuckerman, "Decision-Making by Consensus Doesn't Work in the AI Era," *Harvard Business Review*, April 2026.

² Ethan Mollick, *Co-Existence: The Next Phase of AI* (Portfolio), 20 October 2026. The frontier argument here draws on his essay "Co-Existence and the End of Co-Intelligence," *One Useful Thing*, 4 June 2026.

³ Robert Timmer, *The Paradox of Human-AI Collaborative Decision-Making: Dealing With a Dual-Bounded System*, master's thesis, Nyenrode Business University, 2026.



Better decisions. By design.

ROBERT.TIMMER@THE-PARADOX.COM

THE-PARADOX.COM

THE HAGUE — WORKING GLOBALLY

BEGIN THE INQUIRY →

